

hw3-karim

September 24, 2024

1. Load Advertising.csv dataset using pandas

```
[2332]: import pandas as pd
import sklearn
import numpy as np

df = pd.read_csv("Advertising.csv", index_col=0)
y = df["Sales"]
x = df.drop("Sales", axis=1)
print(x[:5])
print(y[:5])
```

	TV	Radio	Newspaper
1	230.1	37.8	69.2
2	44.5	39.3	45.1
3	17.2	45.9	69.3
4	151.5	41.3	58.5
5	180.8	10.8	58.4

1	22.1
2	10.4
3	9.3
4	18.5
5	12.9

Name: Sales, dtype: float64

2. Standardize each column of the dataset For each predictor x_j , for $j = 0, 1, \dots, j$, compute for the standardized values:

```
[2333]: x_scaled = (x-np.mean(x, axis=0))/np.std(x, axis=0)
print("Scaled X:\n", x_scaled[:5])

x_sk = sklearn.preprocessing.scale(df, axis=0)
print("\nScaled X using sklearn:\n", x_sk[:5])

y_scaled = (y-np.mean(y, axis=0))/np.std(y, axis=0)
print("\nScaled Y:\n", y[:5])
```

Scaled X:

	TV	Radio	Newspaper
1	0.969852	0.981522	1.778945

```

2 -1.197376  1.082808  0.669579
3 -1.516155  1.528463  1.783549
4  0.052050  1.217855  1.286405
5  0.394182 -0.841614  1.281802

```

Scaled X using sklearn:

```

[[ 0.96985227  0.98152247  1.77894547  1.55205313]
 [-1.19737623  1.08280781  0.66957876 -0.69604611]
 [-1.51615499  1.52846331  1.78354865 -0.90740587]
 [ 0.05204968  1.21785493  1.28640506  0.86033029]
 [ 0.3941822  -0.84161366  1.28180188 -0.21568303]]

```

Scaled Y:

```

1    22.1
2    10.4
3     9.3
4    18.5
5    12.9

```

Name: Sales, dtype: float64

3. So you must add an extra column composing of all ones to X.

```

[2334]: x_scaled.insert(0, 'Bias', [1]*len(x))

x_scaled[:5]

```

```

[2334]:   Bias      TV      Radio  Newspaper
1      1  0.969852  0.981522  1.778945
2      1 -1.197376  1.082808  0.669579
3      1 -1.516155  1.528463  1.783549
4      1  0.052050  1.217855  1.286405
5      1  0.394182 -0.841614  1.281802

```

4. Divide the dataset into training and testing, with 0.85 and 0.15 ratio, respectively

```

[2335]: import sklearn.model_selection

test_size=0.15
seed=42
x_train, x_test, y_train, y_test = sklearn.model_selection.
    ↪train_test_split(x_scaled, y_scaled, test_size=test_size, random_state=seed)
y_train[:5]

```

```

[2335]: 10    -0.657617
        19    -0.523115
        56     1.859486
        76    -1.022693
       151     0.399182
Name: Sales, dtype: float64

```

5. Fit the model on the training set. Essentially, you have to optimize the model using the training set, and not including the test set. (Instruction 5 elaborated below)
5. 1. `initialize_weights`: returns a vector `init_w` composing of 4 uniformly distributed numbers between 0 and 1. This serves as the initial weights w_j , for $j = 0, 1, 2, 3$. You can set a random seed so you can objectively assess if your model is working correctly. Seed function is used to save the state of a random function, so that it can generate same random numbers on multiple executions of the code.

```
[2336]: def initialize_weights(use_random=True, seed=42):
    """
    Initializes weights for a linear regression model.

    Parameters:
    - use_random (bool): If True, weights are initialized to random values
    between 0 and 1.
                        If False, weights are initialized to 0.
    - seed (int): Random seed for reproducibility (only used if use_random is
    True).

    Returns:
    - init_w (np.array): A vector of 4 weights.
    """
    if use_random:
        np.random.seed(seed) # Set the random seed for reproducibility
        init_w = np.random.uniform(0, 1, 4) # 4 random numbers between 0 and 1
    else:
        init_w = np.array([0.0, 0.0, 0.0, 0.0]) # Initialize all weights to
    zero

    return init_w
```

5. 2. `predict`: returns a vector of the predicted values \hat{y}_i

```
[2337]: def predict(X, weights):
    """
    Predicts the target values (y_hat) using the input features and weights.

    Parameters:
    - X (np.array): The feature matrix (with a column of ones for the bias
    term).
    - weights (np.array): The vector of weights, including the bias term.

    Returns:
    - y_hat (np.array): The predicted values.
    """
    # Calculate the predicted values using matrix multiplication
    y_hat = X.dot(weights)
```

```
return y_hat
```

5. 3. `compute_cost`: returns a scalar value that tells us how accurate the model is

```
[2338]: def compute_cost(X, y, weights):  
    """  
    Computes the cost function (Mean Squared Error) for the given data and  
    weights.  
  
    Parameters:  
    - X (np.array): The feature matrix (with a column of ones for the bias  
    term).  
    - y (np.array): The actual target values.  
    - weights (np.array): The vector of weights, including the bias term.  
  
    Returns:  
    - cost (float): The cost value (scalar), representing the model's error.  
    """  
    m = len(y) # Number of training examples  
    y_hat = X.dot(weights) # Predicted values  
  
    # Compute the squared differences  
    squared_errors = (y_hat - y) ** 2  
  
    # Compute the cost (Mean Squared Error)  
    cost = (1 / (2 * m)) * np.sum(squared_errors)  
  
    return cost
```

5. 4. `compute_gradient`: returns a matrix `w` that represents the partial derivative of the cost function with respect to with respect to each parameter

```
[2339]: def compute_gradient(X, y, weights):  
    """  
    Computes the gradient of the cost function with respect to the weights.  
  
    Parameters:  
    - X (np.array): The feature matrix (with a column of ones for the bias  
    term).  
    - y (np.array): The actual target values.  
    - weights (np.array): The vector of weights, including the bias term.  
  
    Returns:  
    - gradient (np.array): The gradient matrix (4x1), representing the partial  
    derivatives.  
    """  
    m = len(y) # Number of training examples
```

```

y_hat = X.dot(weights) # Predicted values

# Compute the error (difference between predicted and actual values)
error = y_hat - y

# Compute the gradient for each parameter
gradient = (1 / m) * X.T.dot(error)

return gradient

```

5. 5. update_weights: returns a 4x1 matrix that contains the updated weights.

```

[2340]: def update_weights(weights, gradient, learning_rate=0.01):
        """
        Updates the weights using the gradient and learning rate.

        Parameters:
        - weights (np.array): The current weights (4x1 vector, including the bias
        ↪term).
        - gradient (np.array): The gradient of the cost function with respect to
        ↪the weights (4x1 vector).
        - learning_rate (float): The learning rate (alpha) for gradient descent.

        Returns:
        - updated_weights (np.array): The updated weights after one iteration of
        ↪gradient descent (4x1 vector).
        """
        # Update the weights by subtracting the learning rate times the gradient
        updated_weights = weights - learning_rate * gradient

        return updated_weights

```

5. 6. grad_descent: returns 2 matrices: one matrix for the weights, and one matrix for the cost values per iteration. grad_descent calls the functions 1-5 until the number of iterations is reached.

```

[2341]: def grad_descent(X, y, learning_rate=0.01, iterations=100):
        """
        Performs gradient descent to learn the weights, calling the necessary
        ↪helper functions.

        Parameters:
        - X (np.array): The feature matrix (with a column of ones for the bias
        ↪term).
        - y (np.array): The actual target values.
        - learning_rate (float): The learning rate (alpha) for gradient descent.
        - iterations (int): The number of iterations to run gradient descent.

```

```

Returns:
- weights_history (np.array): A matrix containing the weights at each
↪ iteration.
- cost_history (np.array): A matrix containing the cost at each iteration.
"""
m, n = X.shape # Number of training examples (m) and number of features (n)

# Initialize weights and histories
weights = initialize_weights(seed=42) # Calling the initialize_weights_
↪ function
weights_history = np.zeros((iterations, n))
cost_history = np.zeros(iterations)

for i in range(iterations):
    # Make predictions (y_hat) using the current weights
    y_hat = predict(X, weights) # Call the predict function

    # Compute the gradient
    gradient = compute_gradient(X, y, weights) # Call the compute_gradient_
↪ function

    # Update the weights
    weights = update_weights(weights, gradient, learning_rate) # Call the_
↪ update_weights function

    # Compute the cost using the updated weights
    cost = compute_cost(X, y, weights) # Call the compute_cost function

    # Store the weights and cost for this iteration
    weights_history[i, :] = weights
    cost_history[i] = cost

return weights_history, cost_history

```

5. 7. plot_costs: plot the costs as a function of iteration

```

[2342]: import matplotlib.pyplot as plt

def plot_costs(cost_history):
    """
    Plots the costs as a function of iteration.

    Parameters:
    - cost_history (np.array): The array containing the cost at each iteration.
    """
    plt.figure(figsize=(18, 10))

```

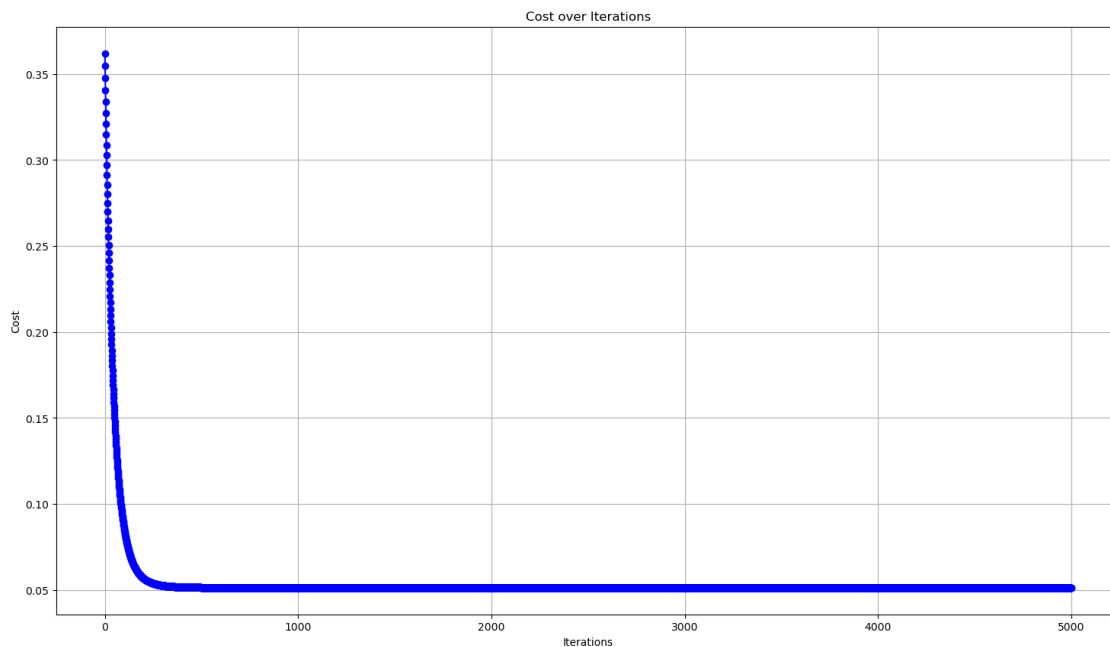
```

plt.plot(range(len(cost_history)), cost_history, color='blue', marker='o',
↪linestyle='-')
plt.title('Cost over Iterations')
plt.xlabel('Iterations')
plt.ylabel('Cost')
plt.grid(True)
plt.show()

learning_rate = 0.01
iterations = 5000
weights_history, cost_history = grad_descent(x_train, y_train, learning_rate,
↪iterations)

# Plot the costs
plot_costs(cost_history)
print("Finalists for the weights")
weights_history[-5:]

```



Finalists for the weights

```

[2342]: array([[4.36682560e-04, 7.37383146e-01, 5.36307180e-01, 3.14254020e-03],
               [4.36682560e-04, 7.37383146e-01, 5.36307180e-01, 3.14254020e-03],
               [4.36682560e-04, 7.37383146e-01, 5.36307180e-01, 3.14254020e-03],
               [4.36682560e-04, 7.37383146e-01, 5.36307180e-01, 3.14254020e-03],
               [4.36682560e-04, 7.37383146e-01, 5.36307180e-01, 3.14254020e-03]])

```

```
[2343]: final_weights = weights_history[iterations-1]
print("Final weights:", final_weights)
```

Final weights: [4.36682560e-04 7.37383146e-01 5.36307180e-01 3.14254020e-03]

5. 8. Predict y for train set and calculate the cost.

```
[2344]: y_train_hat = predict(x_train, final_weights)
print("Predicted y of train set:\n", y_train_hat[:5])

y_train_cost = compute_cost(x_train, y_train, final_weights)
print("\nCost of train set:\n", y_train_cost)
```

Predicted y of train set:

```
10    -0.294970
19    -0.771680
56     1.397679
76    -0.371565
151    0.813103
dtype: float64
```

Cost of train set:

```
0.05133623590933358
```

5. 6. Predict y for test set and calculate the cost

```
[2345]: y_test_hat = predict(x_test, final_weights)
print("Predicted y of test set:\n", y_test_hat[:5])

y_test_cost = compute_cost(x_test, y_test, final_weights)
print("\nCost of test set:\n", y_test_cost)
```

Predicted y of test set:

```
96     0.445526
16     1.304953
31     1.440518
159    -0.668983
129     1.559238
dtype: float64
```

Cost of test set:

```
0.05268626215593916
```

8. Since the data is standardized, you might be surprised that the predictions differ from the original data. In order to revert back a standardized data into the original form, we simply have to equate the previous equation:

```
[2346]: x_mean = np.mean(x, axis=0)
x_std = np.std(x, axis=0)
```



```

scaled_bias = final_weights[0]
scaled_weights = final_weights[1:]

# Adjust the bias/weights for the original scale
bias_orig = scaled_bias - np.sum(scaled_weights * x_mean / x_std)
weights_orig = scaled_weights / x_std

adjusted_weights = np.concatenate([[bias_orig], weights_orig])

# Rescale the weights back to the original scale
print("Weights adjusted for original scale:\n", adjusted_weights)

y_mean = np.mean(y, axis=0)
y_std = np.std(y, axis=0)

x_orig = x
x_orig.insert(0, 'Bias', [1]*len(x))

y_scaled_hat = predict(x_orig, adjusted_weights)
print("\nPredicted y of orig set with adjusted weights:\n", y_scaled_hat[:5])

# Reverse the standardization of y
y_orig_hat = y_scaled_hat * y_std + y_mean

print("\nReversed standardization predicted values:\n", y_orig_hat[:5])

```

Weights adjusted for original scale:

```
[-2.11253587e+00  8.61033279e-03  3.62133702e-02  1.44656825e-04]
```

Predicted y of orig set with adjusted weights:

```

1    1.247577
2   -0.299667
3   -0.292220
4    0.696004
5   -0.156235
dtype: float64

```

Reversed standardization predicted values:

```

1    20.515387
2    12.462916
3    12.501673
4    17.644782
5    13.209389
dtype: float64

```

9. Observe the cost results and analyse.

```
[2347]: m = len(y) # Number of samples
cost = (1 / m) * np.sum((y - y_orig_hat) ** 2)
print("Cost of original unscaled predicted values:\n", cost)
```

Cost of original unscaled predicted values:
2.7919302013509144

1 Questions

1. What are the optimal weights found by your implemented gradient descent? Plug it into the linear model:

$$h_{\theta}(x) = \theta_0 + \theta_1 TV + \theta_2 Radio + \theta_3 Newspaper$$

```
[2348]: print(f"h_0(x)={bias_orig} + {adjusted_weights[1]} * TV + {adjusted_weights[2]} * Radio + {adjusted_weights[3]} * Newspaper")
```

$h_0(x) = -2.112535865333081 + 0.008610332790565228 * TV + 0.036213370182075835 * Radio + 0.0001446568246602006 * Newspaper$

- The sales value is equal to the intercept 3.028... when other variables are held down to 0 (tv=0, radio=0, newspaper=0)

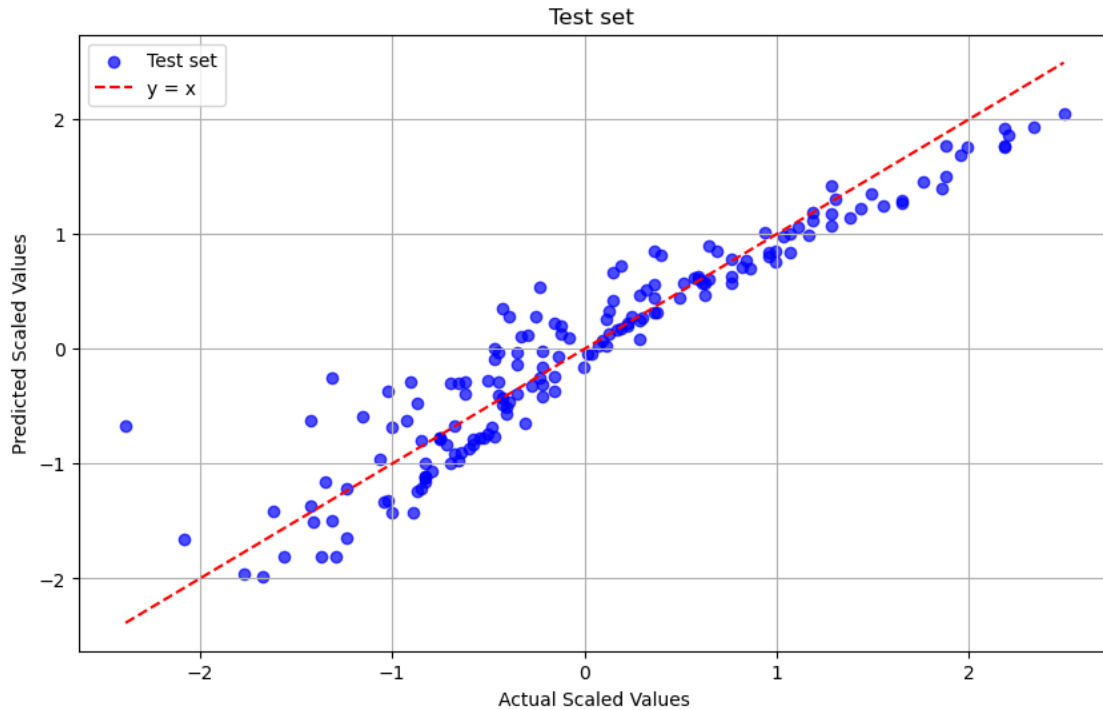
2.

```
[2349]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.scatter(y_train, y_train_hat, color='blue', label='Test set', alpha=0.7)

# Adding a line for reference (y = x)
plt.plot([y_train.min(), y_train.max()],
         [y_train.min(), y_train.max()],
         color='red', linestyle='--', label='y = x')

plt.title('Test set ')
plt.xlabel('Actual Scaled Values')
plt.ylabel('Predicted Scaled Values')
plt.grid(True)
plt.legend()
plt.show()
```




```
[2350]: import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_test_hat, color='blue', label='Test set', alpha=0.7)

# Adding a line for reference (y = x)
plt.plot([y_test.min(), y_test.max()],
         [y_test.min(), y_test.max()],
         color='red', linestyle='--', label='y = x')

plt.title('Test set ')
plt.xlabel('Actual Scaled Values')
plt.ylabel('Predicted Scaled Values')
plt.grid(True)
plt.legend()
plt.show()
```



```
[2351]: import sklearn.metrics as sk_metrics
```

```
print("R2 Score for train:", sk_metrics.r2_score(y_train,y_train_hat))
print("R2 Score for train:", sk_metrics.r2_score(y_test,y_test_hat))
```

R2 Score for train: 0.8936793584593189

R2 Score for train: 0.9110275702091694

3. What happens to the error, r^2 , and cost as the number of iterations increase? Show your data and proof. You can alternatively plot your result data for visualization and check until 50000 iterations or more (actually).

```
[2352]: iteration_list = [100, 500, 1000, 2000, 5000, 10000, 25000, 50000]
```

```
for n_iter in iteration_list:
    weights, cost_history = grad_descent(x_train, y_train, learning_rate,
    ↪n_iter)
    y_train_hat = predict(x_train, weights[-1])
    y_test_hat = predict(x_test, weights[-1])

    # Get the final cost (last value in cost_history)
    final_cost = cost_history[-1]

    # Print the cost and  $R^2$  scores for the current iteration count
```

```

print(f"Iterations: {n_iter}")
print(f"Final Cost: {final_cost}")
print(f"R^2 Score for {n_iter} iterations (Train): {sk_metrics.
↪r2_score(y_train, y_train_hat)}")
print(f"R^2 Score for {n_iter} iterations (Test): {sk_metrics.
↪r2_score(y_test, y_test_hat)}\n")

```

Iterations: 100
 Final Cost: 0.08600217101672748
 R² Score for 100 iterations (Train): 0.8218839804979278
 R² Score for 100 iterations (Test): 0.8098004538502264

Iterations: 500
 Final Cost: 0.05143211521859518
 R² Score for 500 iterations (Train): 0.8934807862521712
 R² Score for 500 iterations (Test): 0.9101235942076645

Iterations: 1000
 Final Cost: 0.05133650395407905
 R² Score for 1000 iterations (Train): 0.8936788033214363
 R² Score for 1000 iterations (Test): 0.9109912615411897

Iterations: 2000
 Final Cost: 0.05133623591159402
 R² Score for 2000 iterations (Train): 0.8936793584546373
 R² Score for 2000 iterations (Test): 0.9110274671841093

Iterations: 5000
 Final Cost: 0.05133623590933358
 R² Score for 5000 iterations (Train): 0.8936793584593189
 R² Score for 5000 iterations (Test): 0.9110275702091694

Iterations: 10000
 Final Cost: 0.0513362359093336
 R² Score for 10000 iterations (Train): 0.8936793584593188
 R² Score for 10000 iterations (Test): 0.9110275702091721

Iterations: 25000
 Final Cost: 0.0513362359093336
 R² Score for 25000 iterations (Train): 0.8936793584593188
 R² Score for 25000 iterations (Test): 0.9110275702091721

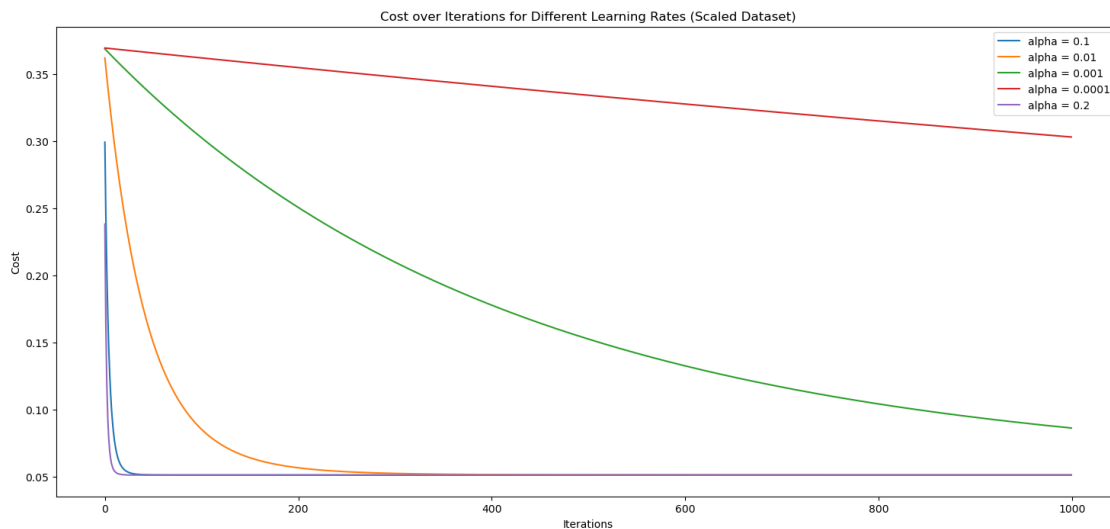
Iterations: 50000
 Final Cost: 0.0513362359093336
 R² Score for 50000 iterations (Train): 0.8936793584593188
 R² Score for 50000 iterations (Test): 0.9110275702091721

- As the iterations increased, the cost rapidly decreased but the difference is plateauing starting from iteration 500. As for the R2 score, it is increasing but the increase is plateauing just like a logarithmic function, starting from 500.
4. Once you determine the optimal number of iterations, check the effect on the cost and error as you change the learning rate. The common learning rates in machine learning include 0.1, 0.01, 0.001, 0.0001, 0.2 but you have the option to include others. Visualize the cost function (vs the optimal number of iterations) of each learning rate in **ONLY ONE PLOT**. Provide your analysis.

```
[2353]: learning_rates = [0.1, 0.01, 0.001, 0.0001, 0.2]
iterations = 1000

plt.figure(figsize=(18, 8))
for alpha in learning_rates:
    weights, cost_history = grad_descent(x_train, y_train, alpha, iterations)
    plt.plot(cost_history, label=f"alpha = {alpha}")

plt.title("Cost over Iterations for Different Learning Rates (Scaled Dataset)")
plt.xlabel("Iterations")
plt.ylabel("Cost")
plt.legend()
plt.show()
```



- Higher learning rates converge rapidly(0.1 and 0.2) while lower learning rates tend to take longer to start converging(0.01 and lower). Learning rates 0.001 and lower won't even show sign of convergence even after the 1000th iteration, taking too long and can waste computing resources.
5. Is there a relationship on the learning rate and the number of iterations?

- Higher learning rates are faster to reach convergence, needing lower iterations. While lower learning rates takes the slow approach, reaching convergence at longer and higher iterations
6. Compare the results with the results of ordinary least squares function

```
[2354]: # Fit the model using OLS
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error

ols_model = LinearRegression()
ols_model.fit(x_train.iloc[:, 1:], y_train) # Removing the bias column
y_train_pred_ols = ols_model.predict(x_train.iloc[:, 1:])
y_test_pred_ols = ols_model.predict(x_test.iloc[:, 1:])

# R2 scores
r2_train = sk_metrics.r2_score(y_train, y_train_hat)
r2_test = sk_metrics.r2_score(y_test, y_test_hat)

# Costs
train_cost = compute_cost(x_train, y_train, adjusted_weights)
test_cost = compute_cost(x_test, y_test, adjusted_weights)

# Compute R2 score for OLS
r2_train_ols = sk_metrics.r2_score(y_train, y_train_pred_ols)
r2_test_ols = sk_metrics.r2_score(y_test, y_test_pred_ols)

# Compute MSE for OLS
mse_train_ols = mean_squared_error(y_train, y_train_pred_ols)
mse_test_ols = mean_squared_error(y_test, y_test_pred_ols)

# Linear regression data
print("-----")
print(f"OLS R^2 score for training set: {r2_train_ols}")
print(f"OLS R^2 score for test set: {r2_test_ols}")
print(f"OLS MSE for training set: {mse_train_ols}")
print(f"OLS MSE for test set: {mse_test_ols}")

print(f"OLS Coefficients: {ols_model.intercept_}, {ols_model.coef_}")

# Compare

print("-----")
print(f"Gradient Descent R^2 score for training set: {r2_train}")
print(f"Gradient Descent R^2 score for test set: {r2_test}")
print(f"Gradient Descent MSE for training set: {train_cost}")
print(f"Gradient Descent MSE for test set: {test_cost}")
```

```
print(f"Gradient Descent Coefficients: {final_weights[0]}, {final_weights[1:]})")
```

```
-----  
OLS R^2 score for training set: 0.8936793584593188  
OLS R^2 score for test set: 0.9110275702091712  
OLS MSE for training set: 0.10267247181866718  
OLS MSE for test set: 0.10537252431187616  
OLS Coefficients: 0.00043668256008587425, [0.73738315 0.53630718 0.00314254]  
-----
```

```
Gradient Descent R^2 score for training set: 0.8936793584593188  
Gradient Descent R^2 score for test set: 0.9110275702091721  
Gradient Descent MSE for training set: 2.7228715674097526  
Gradient Descent MSE for test set: 2.600264391197922  
Gradient Descent Coefficients: 0.0004366825600865185, [0.73738315 0.53630718  
0.00314254]
```

- They are very much identical, everything actually